

Comparative Study of FISTA, ISTA, and BM3D Algorithms for Image Inpainting

Linhongqin
Southwest Petroleum University

Abstract

This paper presents a comprehensive comparative study of image inpainting algorithms, focusing on the Fast Iterative Shrinkage-Thresholding Algorithm (FISTA) and its comparison with the basic Iterative Shrinkage-Thresholding Algorithm (ISTA) and Block-Matching and 3D Filtering (BM3D). Two variants of FISTA are implemented: FISTA-L1 using wavelet-based L1 regularization and FISTA-TV using total variation regularization. Experiments conducted on the Set14 dataset with random pixel missing patterns (60% missing rate) demonstrate that FISTA-L1 achieves the best restoration quality (20.85 dB PSNR), while FISTA-TV offers the fastest computation (0.33 seconds). Convergence analysis confirms FISTA's superior $O(1/k^2)$ convergence rate compared to ISTA's $O(1/k)$. The study provides practical insights for algorithm selection based on specific image characteristics and application requirements.

1 Introduction

Image inpainting, the process of reconstructing missing or corrupted regions in images, is fundamental to computer vision and image processing. Traditional methods range from diffusion-based approaches to patch-based methods and optimization-based techniques. Recently, sparse representation and convex optimization have gained prominence due to their theoretical guarantees and practical effectiveness.

The Iterative Shrinkage-Thresholding Algorithm (ISTA) provides a framework for solving ℓ_1 -regularized linear inverse problems, yet its $O(1/k)$ convergence rate limits practical applications. The Fast Iterative Shrinkage-Thresholding Algorithm (FISTA) accelerates ISTA to achieve $O(1/k^2)$ convergence through Nesterov's momentum technique. Concurrently, BM3D has emerged as a state-of-the-art denoising algorithm leveraging non-local self-similarity.

This study implements and compares these algorithms for image inpainting under random pixel missing patterns. We evaluate two FISTA variants: wavelet-based FISTA-L1 for texture preservation and total variation-based FISTA-TV for edge preservation, providing a systematic performance analysis across quality, speed, and convergence metrics.

2 Methodology

2.1 FISTA Algorithm

FISTA solves optimization problems of the form $\min_x F(x) = f(x) + g(x)$, where f is convex and differentiable (data fidelity term), and g is convex but possibly non-differentiable (regularization term). The algorithm iterates as:

$$y_k = x_k + \frac{t_{k-1} - 1}{t_k}(x_k - x_{k-1}), \quad (1)$$

$$x_{k+1} = \text{prox}_{\lambda g}(y_k - \nabla f(y_k)), \quad (2)$$

$$t_{k+1} = \frac{1 + \sqrt{1 + 4t_k^2}}{2}, \quad (3)$$

where $\text{prox}_{\lambda g}$ denotes the proximal operator.

2.1.1 FISTA-L1 (Wavelet Domain)

For FISTA-L1, we employ wavelet transform \mathcal{W} with L1 regularization: $g(x) = \lambda \|\mathcal{W}(x)\|_1$. The proximal operator corresponds to the soft-thresholding function:

$$\text{prox}_{\lambda g}(z) = \mathcal{W}^{-1}(\text{sign}(\mathcal{W}(z)) \odot \max(|\mathcal{W}(z)| - \lambda, 0)). \quad (4)$$

2.1.2 FISTA-TV (Total Variation)

For FISTA-TV, total variation regularization is used: $g(x) = \lambda \|\nabla x\|_1$, implemented via the Chambolle dual approach for the proximal operator.

2.2 ISTA Algorithm

ISTA serves as the baseline algorithm without acceleration: $x_{k+1} = \text{prox}_{\lambda g}(x_k - \nabla f(x_k))$. We implement ISTA-L1 with wavelet regularization for direct comparison.

2.3 BM3D Algorithm

BM3D employs block-matching to group similar patches, followed by 3D transform-domain filtering and aggregation. For inpainting, we apply iterative BM3D denoising while preserving known pixels.

3 Experimental Setup

3.1 Dataset and Metrics

Experiments are conducted on the Set14 dataset, focusing on the `ppt3.png` image (480×500 pixels). Random pixel missing is simulated with a 60% missing rate. Performance is evaluated using Peak Signal-to-Noise Ratio (PSNR) and Structural Similarity Index (SSIM). PSNR is defined as $\text{PSNR} = 10 \log_{10} (\text{MAX}_I^2 / \text{MSE})$, where $\text{MAX}_I=1.0$ for normalized images and MSE is mean squared error. SSIM measures perceptual similarity considering luminance, contrast, and structure.

3.2 Implementation Details

All algorithms are implemented in Python 3.7. Key parameters include: maximum iterations of 100 for iterative methods; regularization parameters $\lambda = 0.05$ for L1-based methods and $\lambda = 0.01$ for FISTA-TV; Daubechies-4 wavelet with 3-level decomposition; and BM3D with $\sigma_{\text{psd}} = 30$ over 3 iterations.

4 Results and Analysis

4.1 Quantitative and Visual Results

Table 1 presents quantitative comparisons. FISTA-L1 achieves the highest PSNR (20.85 dB) and SSIM (0.5486). FISTA-TV is the fastest (0.33 seconds) but yields lower quality. ISTA achieves comparable quality to FISTA-L1, while BM3D performs poorly for this inpainting task.

Visual comparisons (see Figure 1) show that FISTA-L1 and ISTA produce visually pleasing reconstructions with preserved textures. In contrast, FISTA-TV oversmooths textures but maintains edges well, and BM3D struggles with large missing regions.

4.2 Algorithm Comparison

FISTA demonstrates superior convergence speed compared to ISTA due to Nesterov’s momentum. While both achieve similar final PSNR (20.85 dB vs. 20.82 dB), FISTA reaches near-optimal solutions in fewer iterations, aligning with theoretical expectations of $O(1/k^2)$ versus $O(1/k)$ convergence.

FISTA-L1 outperforms FISTA-TV in PSNR and SSIM for `ppt3.png`, indicating wavelet-based sparse representation better captures texture information. FISTA-TV’s strength in edge preservation makes it suitable for piecewise-smooth images with prominent edges.

Iterative optimization methods (FISTA/ISTA) excel in handling large missing regions by explicitly modeling the data fidelity term. BM3D, designed primarily for additive noise removal, performs poorly with 60% random missing pixels as its block-matching process fails when reference patches are largely missing.

Algorithm	Time (s)	Rel. Speed	Eff. (dB/s)
FISTA-L1	0.98	1.0×	21.3
FISTA-TV	0.33	2.9×	37.2
ISTA	0.88	1.1×	23.7
BM3D	6.82	0.14×	1.4

Table 2: Computational efficiency analysis

4.3 Computational and Convergence Analysis

Table 2 compares computational efficiency. FISTA-TV is 2.9× faster than FISTA-L1 and 20.7× faster than BM3D. The efficiency metric (PSNR/Time) shows FISTA-TV has the highest efficiency (37.2 dB/s), making it suitable for real-time applications.

Convergence analysis (Figure 2) illustrates that FISTA algorithms show rapid initial loss reduction, with FISTA-L1 achieving 54.2% loss reduction over 100 iterations. ISTA shows similar final loss but slower convergence. The PSNR evolution demonstrates FISTA reaches near-optimal PSNR within 50 iterations, while ISTA requires more iterations. The convergence characteristics are threefold: (1) FISTA exhibits rapid initial convergence, reaching 90% of final PSNR within 30 iterations; (2) ISTA shows steady but slower convergence, requiring approximately 70 iterations; (3) BM3D operates as a non-iterative, single-pass algorithm.




Figure 2: Convergence behavior: (a) Loss function vs. iterations, (b) PSNR evolution during iterations. FISTA shows faster convergence than ISTA.

5 Discussion

5.1 Limitations and Extensions

Several limitations warrant further investigation. First, performance depends on regularization parameters; automatic parameter selection would enhance practicality. Second, the current implementation focuses on grayscale images; extension to color images requires channel-wise or vector-valued regularization. Third, random missing pixels represent one scenario; structured patterns (e.g., text removal) may require different approaches. Finally, integration with deep learning could combine optimization-based methods with learned priors.

Algorithm	PSNR (dB)	SSIM	Time (s)	Iterations
Corrupted Image	9.40	—	—	—
FISTA-L1	20.85	0.5486	0.98	100
FISTA-TV	12.26	0.1620	0.33	100
ISTA	20.82	0.5476	0.88	100
BM3D	9.46	0.1218	6.82	—

Table 1: Quantitative results on ppt3.png (480×500, 60% missing)




Figure 1: Visual comparison of inpainting results on ppt3.png. Top row: original, corrupted (PSNR: 9.40 dB), and mask. Middle row: restored images by FISTA-L1, FISTA-TV, ISTA, and BM3D. Bottom row: convergence curves, PSNR evolution, performance comparison, and error map.

5.2 Application Recommendations

Based on our findings, we recommend FISTA-TV for real-time applications due to its speed advantage, and FISTA-L1 for high-quality restoration requiring optimal PSNR/SSIM. Wavelet-based methods (FISTA-L1/ISTA) are suitable for texture-rich images, while TV-based methods (FISTA-TV) excel with edge-prominent images. BM3D remains effective for small noise removal but not for extensive inpainting tasks.

6 Conclusion

This study presents a comprehensive comparison of FISTA, ISTA, and BM3D for image inpainting. The key findings are: (1) FISTA-L1 achieves the best restoration quality (20.85 dB PSNR) on the Set14

dataset with 60% random pixel missing; (2) FISTA-TV offers the fastest computation (0.33 seconds), suitable for time-sensitive applications; (3) FISTA demonstrates superior convergence speed compared to ISTA, validating the theoretical $O(1/k^2)$ advantage; (4) BM3D performs poorly for large missing regions; (5) the choice between wavelet and TV regularization depends on image characteristics. Future work includes adaptive parameter selection, extension to video inpainting, and integration with deep learning.

Appendix: Project Repository

The complete source code, experimental results, and this report are publicly available at: https://github.com/linhongqin123/FISTA_Project.git.